

Asking Questions in Deep Learning

.....

Mathematical Approaches to Neural Networks

Tyler Farghly

Introduction

Google AI listens to 15 years of sea-bottom recordings for hidden whale songs

Devin Coldewey @techcrunch / 3 weeks ago

Comment



Artificial Intelligence 'outsmarts cancer'

By James Gallagher
Health editor, BBC News website

8 June 2016

f t e Share




Portrait by AI program sells for \$432,000

25 October 2018

f t e Share



Introduction



Introduction



Introduction



Artificial Intelligence: it
will kill us | Jay Tuck |
TEDxHamburgSalon

TEDx Talks


YouTube - 31 Jan 2017

Introduction

'It's going to create a revolution': how AI is transforming the NHS

Technology is making impressive inroads into cancer treatment, saving lives and money

●



▲ Time-consuming tasks such as delineating tumours can be done in minutes with AI Photograph: FirstSignal/Getty Images/Stockphoto




How can we fix this?

Introduction

How can we fix this?
Maths!


What is a Neural Network

The Deep Neural Network




What is a Neural Network

The Deep Neural Network




What is a Neural Network

The Deep Neural Network




What is a Neural Network

The Perceptron







What is a Neural Network

The Perceptron




What is a Neural Network

Activation Functions



What is a Neural Network

The Perceptron



What is a Neural Network

Multilayer Perceptron Networks

$$\begin{pmatrix} w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \\ w_{31} & w_{32} & w_{33} \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \end{pmatrix} + \begin{pmatrix} b_1 \\ b_2 \\ b_3 \end{pmatrix} = \begin{pmatrix} w_{11}x_1 + w_{12}x_2 + w_{13}x_3 + b_1 \\ w_{21}x_1 + w_{22}x_2 + w_{23}x_3 + b_2 \\ w_{31}x_1 + w_{32}x_2 + w_{33}x_3 + b_3 \end{pmatrix}$$

What is a Neural Network


Multilayer Perceptron Networks

$$\begin{pmatrix} w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \\ w_{31} & w_{32} & w_{33} \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \end{pmatrix} + \begin{pmatrix} b_1 \\ b_2 \\ b_3 \end{pmatrix} = \begin{pmatrix} w_{11}x_1 + w_{12}x_2 + w_{13}x_3 + b_1 \\ w_{21}x_1 + w_{22}x_2 + w_{23}x_3 + b_2 \\ w_{31}x_1 + w_{32}x_2 + w_{33}x_3 + b_3 \end{pmatrix}$$

$$\mathbf{W}\mathbf{x} + \mathbf{b}$$

What is a Neural Network


Multilayer Perceptron Networks



$$\phi(\mathbf{W}\mathbf{x} + \mathbf{b})$$

What is a Neural Network


Multilayer Perceptron Networks



$$\phi(\mathbf{W}_1\mathbf{x} + \mathbf{b}_1)$$

What is a Neural Network

Multilayer Perceptron Networks



$$\phi\left(\mathbf{W}_2 \phi(\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1) + \mathbf{b}_2\right)$$

What is a Neural Network

Multilayer Perceptron Networks



$$l_i(\mathbf{v}) = \phi(\mathbf{W}_i \mathbf{v} + \mathbf{b}_i)$$

$$N(\mathbf{x}) = l_K \circ \dots \circ l_1(\mathbf{x})$$

Let \mathcal{W} denote the set of all weights, \mathcal{B} denote the set of all biases in the NN

What is a Neural Network

Example



What is a Neural Network


Training Data

We denote training data by


$$\mathcal{D} = \{(\mathbf{x}_i, \mathbf{y}_i)\}_{i=1}^d$$

Example:

$$\mathbf{x}_1 = \begin{pmatrix} 0.7 \\ \vdots \\ 0.1 \end{pmatrix} \rightarrow$$



$$\mathbf{x}_2 = \begin{pmatrix} 0.4 \\ \vdots \\ 0.9 \end{pmatrix} \rightarrow$$



What is a Neural Network


Training Data

We denote training data by

$$\mathcal{D} = \{(\mathbf{x}_i, \mathbf{y}_i)\}_{i=1}^d$$


Example:

$$\mathbf{x}_1 = \begin{pmatrix} 0.7 \\ \vdots \\ 0.1 \end{pmatrix} \rightarrow$$



$$\mathbf{y}_1 = 1$$

$$\mathbf{x}_2 = \begin{pmatrix} 0.4 \\ \vdots \\ 0.9 \end{pmatrix} \rightarrow$$



$$\mathbf{y}_2 = 0$$

What is a Neural Network

Training and Evaluation

Given data $\mathcal{D} = \{(\mathbf{x}_i, \mathbf{y}_i)\}_{i=1}^d$


$$L(N, \mathcal{D}) = \frac{1}{d} \sum_{i=1}^d \|\mathbf{y}_i - N(\mathbf{x}_i)\|^2 \quad - \text{mean squared error}$$

What is a Neural Network

Training and Evaluation

Aim to find \mathcal{W} , \mathcal{B} such that L is minimised.

Achieve this via SGD and backpropagation.



Schematic of gradient descent.

What is a Neural Network

The Complete Process

- We have a NN with parameters \mathcal{W} and \mathcal{B}
- Adjust these parameters such that L is minimum given our data
- Use this new NN to obtain our prediction given \mathbf{x} , $N(\mathbf{x})$

What is a Neural Network


Black Box Problem



“A black box is a device, system or object which can be viewed in terms of its inputs and outputs, without any knowledge of its internal workings”

What is a Neural Network

So why are they so mysterious?




What is a Neural Network

So why are they rubbish?

$$\begin{aligned} N(x) = & \phi(w\phi(w\phi(w\phi(wx + b) + w\phi(wx + b) + w\phi(wx + b) + b) \\ & + w\phi(w\phi(wx + b) + w\phi(wx + b) + w\phi(wx + b) + b) + w\phi(w\phi(wx + \\ & w\phi(wx + b) + w\phi(wx + b) + b) + b) + w\phi(w\phi(w\phi(wx + b) + w\phi(wx \\ & + b) + w\phi(wx + b) + b) + w\phi(w\phi(wx + b) + w\phi(wx + b) + w\phi(wx \\ & + b) + b) + w\phi(w\phi(wx + b) + w\phi(wx + b) + w\phi(wx + b) + b) + b) \\ & + w\phi(w\phi(w\phi(wx + b) + w\phi(wx + b) + w\phi(wx + b) + b) + \\ & w\phi(w\phi(wx + b) + w\phi(wx + b) + w\phi(wx + b) + b) + w\phi(w\phi(wx + b) \\ & + w\phi(wx + b) + w\phi(wx + b) + b) + b) + b) \end{aligned}$$

What is a Neural Network

So why are they rubbish?



What is a Neural Network

Why are they great?

Theorem 2.3.1 (universal approximation theorem):

Let $\varphi(\cdot)$ be an arbitrary activation function. Let $X \subseteq \mathbb{R}^m$ and X is compact. The space of continuous functions on X is denoted by $C(X)$. Then $\forall f \in C(X), \forall \varepsilon > 0: \exists n \in \mathbb{N}, a_{ij}, b_i, w_i \in \mathbb{R}, i \in \{1 \dots n\}, j \in \{1 \dots m\}$:

$$(A_n f)(x_1, \dots, x_m) = \sum_{i=1}^n w_i \varphi\left(\sum_{j=1}^m a_{ij} x_j + b_i\right)$$

as an approximation of the function $f(\cdot)$; that is

$$\|f - A_n f\| < \varepsilon$$

The Learning Process

Introduction

$$\frac{1}{d} \sum_{i=1}^d (Y_i - N(X_i))^2 \rightarrow E \left[(Y - N(X))^2 \right]$$

as $d \rightarrow \infty$

The Learning Process

Bias/Variance Tradeoff

$$\mathcal{D} \rightarrow \mathcal{W}, \mathcal{B} \rightarrow N(\mathbf{x})$$

The Learning Process

Bias/Variance Tradeoff

$$\mathcal{D} \rightarrow \mathcal{W}, \mathcal{B} \rightarrow N(\mathbf{x})$$

$$E_{\mathcal{D}} \left(E \left[(y - N(\mathbf{x}))^2 \right] \right) = E \left[(y - E_{\mathcal{D}}[N(\mathbf{x})])^2 + Var_{\mathcal{D}}(N(\mathbf{x})) \right]$$

The Learning Process


Bias/Variance Tradeoff

$$\mathcal{D} \rightarrow \mathcal{W}, \mathcal{B} \rightarrow N(\mathbf{x})$$

$$E_{\mathcal{D}} \left(E \left[(y - N(\mathbf{x}))^2 \right] \right) = E \left[\underbrace{(y - E_{\mathcal{D}}[N(\mathbf{x})])^2}_{\text{squared bias}} + \underbrace{Var_{\mathcal{D}}(N(\mathbf{x}))}_{\text{variance}} \right]$$


The Learning Process

Bias/Variance Tradeoff



The Learning Process

Bias/Variance Tradeoff



The Learning Process


Example: Classification

1




0.93

0




0.01

0



0.08

1



0.96

The Learning Process


Example: Classification

1



0.93

0




0.01

0




0.08

1



0.96

0



0.42

The Learning Process

Example: Classification

Split the input space into volumes $V(v)$:

$V(v)$ – volume of size dv and centered at a point v

The Learning Process

Example: Classification

Split the input space into volumes $V(v)$:

$V(v)$ – volume of size dv and centered at a point v

By continuity, $N(v) = N(x) \forall x \in V(v)$

The Learning Process

Example: Classification

Let $p(v) := p(y = 1 | x \in V(v))$

$$E_{V(v)} \left[(y - N(x))^2 \right] = [1 - N(v)]^2 p(v) + N(v)^2 [1 - p(v)]$$

The Learning Process

Example: Classification

$$E \left[(y - N(x))^2 \right] = \sum_V \left[(1 - N(v))p(v) + N(v)^2(1 - p(v)) \right] dv$$

The Learning Process

Example: Classification

$$E \left[(y - N(x))^2 \right] = \sum_V \left[(1 - N(v))p(v) + N(v)^2(1 - p(v)) \right] dv$$

Differentiating each term w.r.t. $N(v)$ and setting to 0 gives

$$-2p(v)[1 - N(v)] + 2(1 - p(v))N(v) = 0$$

From this we deduce $p(v) = N(v)$ minimises the error in any $V(v)$.

The Learning Process

Generalisation Error

Suppose we are creating a classifier, i.e. $y \in \{0, 1\}$.

The VC dimension, h , is the largest number of data points we can always expect the NN to fit to.

The Learning Process

Generalisation Error

Suppose we are creating a classifier, i.e. $y \in \{0, 1\}$.

The VC dimension, h , is the largest number of data points we can always expect the NN to fit to.

ν = average rate of error

ν_{emp} = rate of error on training data

The Learning Process

Generalisation Error

If $\nu(\mathcal{W}, \mathcal{B})$ and $\nu_{emp}(\mathcal{W}, \mathcal{B})$ are sufficiently close to each other with probability α then

$$\nu(\mathcal{W}, \mathcal{B}) < \nu_{emp}(\mathcal{W}, \mathcal{B}) + \varepsilon_1(h, d, \alpha)$$

The Learning Process

Generalisation Error

If $\nu(\mathcal{W}, \mathcal{B})$ and $\nu_{emp}(\mathcal{W}, \mathcal{B})$ are sufficiently close to each other with probability α then


$$\nu(\mathcal{W}, \mathcal{B}) < \nu_{emp}(\mathcal{W}, \mathcal{B}) + \varepsilon_1(h, d, \alpha)$$

$$\varepsilon_1(h, d, \alpha) = 2\varepsilon_0^2 \left(1 + \sqrt{1 + \frac{\nu_{emp}(\mathcal{W}, \mathcal{B})}{\varepsilon_0^2}} \right)$$

$$\varepsilon_0 = \sqrt{\frac{h}{d} \left[\ln \frac{2d}{h} + 1 \right] - \frac{1}{d} \ln \alpha}$$

The Learning Process

Generalisation Error



The Learning Process

Generalisation Error

$\nu_{gene}(\mathcal{W}, \mathcal{B}) = \text{rate of error on new data}$

$$P(\nu_{gene}(\mathcal{W}, \mathcal{B}) < \nu_{train}(\mathcal{W}, \mathcal{B}) + \varepsilon_1) > 1 - \alpha$$

Bayesian Neural Networks


Introduction

1



0.93

0



0.01

Bayesian Neural Networks


Introduction

1



0.93

0



0.01



0.78

Bayesian Neural Networks


Introduction

1



0.93

0



0.01



0.78

No notion of uncertainty!

Bayesian Neural Networks

Introduction

“Bayesian learning is distinguished by its use of probability to express all forms of uncertainty”[RADFORD]

$$\mathcal{W} \quad \mathcal{B}$$

Bayesian Neural Networks

Introduction

$$p(\mathcal{W}, \mathcal{B})$$


Bayesian Neural Networks

Introduction

$$p(\mathcal{W}, \mathcal{B}) \rightarrow p(\mathcal{W}, \mathcal{B} | \mathcal{D})$$

Bayesian Neural Networks

Example



Bayesian Neural Networks

Posterior Distribution

$$p(\mathcal{W}, \mathcal{B} | \mathcal{D}) = \frac{p(\mathcal{D} | \mathcal{W}, \mathcal{B}) p(\mathcal{W}, \mathcal{B})}{\int p(\mathcal{D} | \mathcal{W}', \mathcal{B}') d\mathcal{W}' d\mathcal{B}'}$$

$$\propto p(\mathcal{D} | \mathcal{W}, \mathcal{B}) p(\mathcal{W}, \mathcal{B})$$

Bayesian Neural Networks

Uncertainty

Suppose \mathbf{y}_{pred} is the prediction obtained from the BNN.

We can evaluate the uncertainty of such a prediction by analysing the distribution of $\mathbf{y}|\mathbf{x}, \mathcal{D}$ around $\mathbf{y} = \mathbf{y}_{pred}$

Bayesian Neural Networks

Uncertainty


Suppose \mathbf{y}_{pred} is the prediction obtained from the BNN.

We can evaluate the uncertainty of such a prediction by analysing the distribution of $\mathbf{y}|\mathbf{x}, \mathcal{D}$ around $\mathbf{y} = \mathbf{y}_{pred}$

$$p(\mathbf{y}|\mathbf{x}, \mathcal{D}) = \int p(\mathbf{y}|\mathbf{x}, \mathcal{W}, \mathcal{B}) p(\mathcal{W}, \mathcal{B}|\mathcal{D}) d\mathcal{W} d\mathcal{B}$$


Bayesian Neural Networks

Uncertainty



Bayesian Neural Networks

Uncertainty



Bayesian Neural Networks

Traditional approach

- We have a NN with parameters \mathcal{W} and \mathcal{B}
- Adjust these parameters such that L is minimum given our data
- Use this new NN to obtain our prediction given \mathbf{x} , $N(\mathbf{x})$

Bayesian Neural Networks

Traditional approach

- We have a NN with parameters \mathcal{W} and \mathcal{B}
- Adjust these parameters such that L is minimum given our data
- Use this new NN to obtain our prediction given \mathbf{x} , $N(\mathbf{x})$

Bayesian approach

- Set the prior distributions and sample from $\mathcal{W}, \mathcal{B} | \mathcal{D}$
- Average value of $N(\mathbf{x})$ over the sampled parameters is \mathbf{y}_{pred}
- Evaluate distribution of $\mathbf{y} | \mathbf{x}, \mathcal{D}$ to evaluate uncertainty

Bayesian Neural Networks

Prior Distribution


$$w, b \sim \text{Normal} \left(0, \sigma^2 \right) \quad \text{for all } w \in \mathcal{W}, b \in \mathcal{B}$$

$$p(\mathcal{W}, B) = \prod p(w) \prod p(b)$$

$$\begin{aligned} -\log p(\mathcal{W}, B) &= \sum \log p(w) + \sum \log p(b) \\ &\propto \frac{1}{2\sigma^2} (|\mathcal{W}|^2 + |\mathcal{B}|^2) \end{aligned}$$

Bayesian Neural Networks

Infinite Networks



Bayesian Neural Networks

Infinite Networks

$$w_i, v_i \sim N(0, H^{-1} \sigma_w^2) \quad b_i, c \sim N(0, \sigma_b^2)$$

$$N(x) = c + \sum_{i=1}^H v_i h_i(x)$$

$$h_i(x) = \phi(w_i x + b_i)$$

Bayesian Neural Networks

By CLT,

$$N(\mathbf{x}) \xrightarrow{\mathcal{D}} \text{Normal}\left(0, \sigma_b^2 + \omega_w^2 V(x)\right) \quad \text{as } H \rightarrow \infty$$

Bayesian Neural Networks

Infinite Networks


$$E(N(x_p), N(x_q)) = \sigma_b^2 + \sigma_w^2 C(x_p, x_q)$$

where $C(x_p, x_q) = E(h_i(x_p)h_i(x_q))$

Distributions over functions of this sort are known as
Gaussian processes


Bayesian Neural Networks

Gaussian Process



Bayesian Neural Networks

Gaussian Process



Smooth, flexible functions! No overfitting as a result of increased flexibility

Bayesian Neural Networks

Pro's and Con's

Pro's

- Can evaluate uncertainty
- Don't need to worry about overfitting
- Can apply well-studied ideas in statistics

Bayesian Neural Networks

Pro's and Con's

Pro's

- Can evaluate uncertainty
- Don't need to worry about overfitting
- Can apply well-studied ideas in statistics

Con's

- Very computationally expensive
- Difficult to establish a prior distribution
- Occasionally results rely heavily on choice of prior

Information Theory in Neural Networks

Dropout Regularization


Introduction

A weird method to reduce overfitting

Dropout Regularization

Introduction


Recall SGD as our method for minimising L :




Schematic of gradient descent.

Dropout Regularization

Introduction




(a) Standard Neural Net



(b) After applying dropout.

Dropout Regularization

Introduction



Dropout Regularization

First Paper

Improving neural networks by preventing co-adaptation of feature detectors

G. E. Hinton*, N. Srivastava, A. Krizhevsky, I. Sutskever and R. R. Salakhutdinov

Department of Computer Science, University of Toronto,
6 King's College Rd, Toronto, Ontario M5S 3G4, Canada

Dropout Regularization

First Paper

Improving neural networks by preventing co-adaptation of feature detectors

G. E. Hinton*, N. Srivastava, A. Krizhevsky, I. Sutskever and R. R. Salakhutdinov

Department of Computer Science, University of Toronto,
6 King's College Rd, Toronto, Ontario M5S 3G4, Canada

“Overfitting can be reduced by using dropout to prevent complex
co-adaptations on the training data.”

Dropout Regularization

Formulation

If layer i is a dropout layer

$$l_i(\mathbf{x}) = \frac{1}{p} \phi(W\mathbf{x} + \mathbf{b}) \circ \mathbf{r}_i$$

$$(\mathbf{r}_i)_j \sim \text{Bernoulli}(p_i)$$

Dropout Regularization

Suprising Property of Dropout

Suppose we have the data set $\{(\mathbf{0}, \mathbf{0}), (\mathbf{1}, \mathbf{1})\}$ and we are training a neural network of linear activation where all biases are 0 i.e.

$$N(\mathbf{x}) = W_K \dots W_1 \mathbf{x}$$

Dropout Regularization

Suprising Property of Dropout

Suppose we have the data set $\{(\mathbf{0}, \mathbf{0}), (\mathbf{1}, \mathbf{1})\}$ and we are training a neural network of linear activation where all biases are 0 i.e.

$$N(\mathbf{x}) = W_K \dots W_1 \mathbf{x}$$

If we insert dropout into every layer, for the NN to be optimum there must be **at least one negative weight!**

What's Next

What's Next

Who knows?

Questions

<https://bit.ly/2QV9vjg>

<https://bit.ly/2RORWyq>